





Building an Openstack cluster without clustering

Egon Rijpkema & Wim Nap
Server & Data Infrastructure
Team HPC



openstack®
CLOUD SOFTWARE



docker



- › Integration in Federated Cloud
- › Network modifiable (Opencontrail)
- › Automated installation
- › Easily upgradable
- › Access to hardware



- › Dockers
 - › No Clustering software
 - › Emulating the openstack [install guide](#)
 - › [Dockerfile](#) example
 - › Images built using jenkins
 - › Hosted on vmware Harbor registry



- › Ansible
 - › Powerful
 - › Shallow learning curve
 - › imperative
 - › [Playbooks](#)





- › Systemd unit files
 - › Robust
 - › Integrated logging
 - › [Unit file](#) example



- › openstack components (11)
- › Mariadb Galera cluster
- › Rabbitmq cluster



- › Docker clustering
- › Service discovery (consul)
- › Overlay network



- › We looked at existing solutions.
 - › Not in docker (openstack ansible)
 - › no opencontrail integration
 - › ill supported
 - › .. all this ~~might~~ will change



- › Orchestration (zookeeper, consul etc)
- › dynamic scheduling
- › SDN
- › Dynamic resource allocation (service)
- › Kubernetes on top of openstack
- › Or....?



Questions?

